

Practice5

September 9, 2025

pandas adopts significant parts of NumPy's idiomatic style of array-based computing, especially array-based functions and a preference for data processing without for loops

While pandas adopts many coding idioms from NumPy, the biggest difference is that pandas is designed for working with tabular or heterogeneous data. NumPy, by contrast, is best suited for working with homogeneous numerical array data.

```
[1]: from pandas import Series, DataFrame
```

A Series is a one-dimensional array-like object containing a sequence of values (of similar types to NumPy types) and an associated array of data labels, called its index. The simplest Series is formed from only an array of data.

```
[2]: import pandas as pd
pd.Series([1,1.5,'a',[1,2,3]])
```

```
[2]: 0          1
1          1.5
2          a
3    [1, 2, 3]
dtype: object
```

You can get the array representation and index object of the Series via its values and index attributes, respectively

```
[3]: series=pd.Series([1,1.5,'a',[1,2,3]])
series.index
```

```
[3]: RangeIndex(start=0, stop=4, step=1)
```

```
[4]: series.values
```

```
[4]: array([1, 1.5, 'a', list([1, 2, 3])], dtype=object)
```

```
[5]: list([1, 2, 3])
```

```
[5]: [1, 2, 3]
```

```
[6]: ser=pd.Series(['ankit','summi','kiioo'], index=['a','b','c'])
```

```
[7]: ser
```

```
[7]: a    ankit  
     b    summi  
     c    kiioo  
     dtype: object
```

```
[8]: ser.iloc[0]  #accessing first element using index
```

```
[8]: 'ankit'
```

```
[9]: ser.loc['a']  #accessing first element using label
```

```
[9]: 'ankit'
```

```
[10]: ser['a']
```

```
[10]: 'ankit'
```

```
[11]: ser[['a','c']]  #accessing multiple elements using label
```

```
[11]: a    ankit  
     c    kiioo  
     dtype: object
```

```
[12]: ser[1]
```

```
/tmp/ipykernel_5585/4267038266.py:1: FutureWarning: Series.__getitem__ treating  
keys as positions is deprecated. In a future version, integer keys will always  
be treated as labels (consistent with DataFrame behavior). To access a value by  
position, use `ser.iloc[pos]`  
    ser[1]
```

```
[12]: 'summi'
```

```
[13]: ser
```

```
[13]: a    ankit  
     b    summi  
     c    kiioo  
     dtype: object
```

```
[14]: ser['d'] = 'soumi'  #adding new elements  
     ser
```

```
[14]: a    ankit  
     b    summi
```

```
c    kiioo
d    soumi
dtype: object
```

Using NumPy functions or NumPy-like operations, such as filtering with a boolean array, scalar multiplication, or applying math functions, will preserve the index-value link

```
[15]: ser[(ser=='soumi') | (ser=='kiioo')]
```

```
[15]: c    kiioo
      d    soumi
      dtype: object
```

```
[16]: ser*2
```

```
[16]: a    ankitankit
      b    summisummi
      c    kiiookiioo
      d    soumisoumi
      dtype: object
```

```
[17]: ser+ser
```

```
[17]: a    ankitankit
      b    summisummi
      c    kiiookiioo
      d    soumisoumi
      dtype: object
```

```
[18]: ser+'3'
```

```
[18]: a    ankit3
      b    summi3
      c    kiioo3
      d    soumi3
      dtype: object
```

Another way to think about a Series is as a fixed-length, ordered dict, as it is a mapping of index values to data values. It can be used in many contexts where you might use a dict

```
[19]: 'b' in ser
```

```
[19]: True
```

```
[20]: 'f' in ser
```

```
[20]: False
```

```
[21]: ser1=pd.Series({'a':'ankkit','k':'kiioo'})
ser1
```

```
[21]: a    ankkit
      k    kiioo
      dtype: object
```

I will use the terms “missing” or “NA” interchangeably to refer to missing data. The `isnull` and `notnull` functions in pandas should be used to detect missing data

```
[22]: obj=pd.Series([1,2,None,4,5],index=['a','b','c','d','e'])
      print(obj,'\n\n',obj.isnull(),'\n\n',obj.isnull().sum())
```

```
a    1.0
b    2.0
c    NaN
d    4.0
e    5.0
dtype: float64
```

```
      a    False
      b    False
      c     True
      d    False
      e    False
dtype: bool
```

```
1
```

Both the Series object itself and its index have a `name` attribute

```
[23]: obj.name='My Object'
      obj.index.name = 'My Index'
      obj
```

```
[23]: My Index
      a    1.0
      b    2.0
      c    NaN
      d    4.0
      e    5.0
      Name: My Object, dtype: float64
```

```
[24]: import numpy as np
      obj=pd.Series([1,2,3,4,5],index=['a','b','c','d','e'],dtype=np.int16)
      obj
```

```
[24]: a    1
      b    2
      c    3
      d    4
      e    5
      dtype: int16
```

Dataframe

A DataFrame represents a rectangular table of data and contains an ordered collection of columns, each of which can be a different value type (numeric, string, boolean, etc.).

There are many ways to construct a DataFrame, though one of the most common is from a dict of equal-length lists or NumPy arrays

```
[25]: data={'name':['ankit','kiio','summi','soumi'],'institute':
      ↪['ISI','IIIT-D','ECIL','IISc'],'Priority':[4,1,2,3],'Marks':[0,100.0,100.
      ↪0,100.0]}
      df=pd.
      ↪DataFrame(data,index=['a','b','c','d'],columns=['name','institute','Marks','Priority'])
      df
```

```
[25]:   name institute  Marks  Priority
a  ankit        ISI    0.0         4
b  kiio     IIIT-D  100.0         1
c  summi        ECIL  100.0         2
d  soumi        IISc  100.0         3
```

```
[26]: df.head(3)
```

```
[26]:   name institute  Marks  Priority
a  ankit        ISI    0.0         4
b  kiio     IIIT-D  100.0         1
c  summi        ECIL  100.0         2
```

```
[27]: df.columns
```

```
[27]: Index(['name', 'institute', 'Marks', 'Priority'], dtype='object')
```

```
[28]: df['name'] #dict like notation
```

```
[28]: a    ankit
      b    kiio
      c    summi
      d    soumi
      Name: name, dtype: object
```

```
[29]: df.name
```

```
[29]: a    ankit
      b    kiio
      c    summi
      d    soumi
      Name: name, dtype: object
```

```
[30]: # column to numpy array
      df.name.values
```

```
[30]: array(['ankit', 'kiio', 'summi', 'soumi'], dtype=object)
```

```
[31]: df['name'].values
```

```
[31]: array(['ankit', 'kiio', 'summi', 'soumi'], dtype=object)
```

```
[32]: x=pd.Series({'a':1,'b':[1,2]},index=['b','a'])
```

```
[33]: x
```

```
[33]: b    [1, 2]
      a         1
      dtype: object
```

```
[34]: x['b']
```

```
[34]: [1, 2]
```

```
[35]: df.loc[['c','d']]
```

```
[35]:   name institute  Marks  Priority
c  summi      ECIL  100.0         2
d  soumi      IISc  100.0         3
```

```
[36]: df.iloc[[0,1]]
```

```
[36]:   name institute  Marks  Priority
a  ankit      ISI    0.0         4
b  kiio    IIIT-D  100.0         1
```

```
[37]: df.loc[df['name']=='kiio']
```

```
[37]:   name institute  Marks  Priority
b  kiio    IIIT-D  100.0         1
```

```
[38]: df['Marrital Statu']='No'
```

```
[39]: df
```

```
[39]:
```

	name	institute	Marks	Priority	Marrital	Statu
a	ankit	ISI	0.0	4		No
b	kiio	IIIT-D	100.0	1		No
c	summi	ECIL	100.0	2		No
d	soumi	IISc	100.0	3		No

```
[40]: df['rank']=np.arange(1,5,1)
print(df)
```

	name	institute	Marks	Priority	Marrital	Statu	rank
a	ankit	ISI	0.0	4		No	1
b	kiio	IIIT-D	100.0	1		No	2
c	summi	ECIL	100.0	2		No	3
d	soumi	IISc	100.0	3		No	4

```
[41]: df['rank']=pd.Series([4,1,2,3],index=['a','b','c','d'])
print(df)
```

	name	institute	Marks	Priority	Marrital	Statu	rank
a	ankit	ISI	0.0	4		No	4
b	kiio	IIIT-D	100.0	1		No	1
c	summi	ECIL	100.0	2		No	2
d	soumi	IISc	100.0	3		No	3

```
[42]: # You can transpose the DataFrame (swap rows and columns) with similar syntax
↳ to a NumPy array
df.T
```

```
[42]:
```

	a	b	c	d
name	ankit	kiio	summi	soumi
institute	ISI	IIIT-D	ECIL	IISc
Marks	0.0	100.0	100.0	100.0
Priority	4	1	2	3
Marrital Statu	No	No	No	No
rank	4	1	2	3

```
[43]: df.transpose()
```

```
[43]:
```

	a	b	c	d
name	ankit	kiio	summi	soumi
institute	ISI	IIIT-D	ECIL	IISc
Marks	0.0	100.0	100.0	100.0
Priority	4	1	2	3
Marrital Statu	No	No	No	No
rank	4	1	2	3

a DataFrame's index and columns have their name attributes like series.

```
[44]: print(df['rank'].values)
      print(df.values)
```

```
[4 1 2 3]
[['ankit' 'ISI' 0.0 4 'No' 4]
 ['kiio' 'IIIT-D' 100.0 1 'No' 1]
 ['summi' 'ECIL' 100.0 2 'No' 2]
 ['soumi' 'IISc' 100.0 3 'No' 3]]
```

```
[45]: df1=pd.DataFrame({'Name':['ankit','kiio'],'Age':[30,25]},index=['a','b'])
      new_index=df1.index
      df2=pd.DataFrame({'College':['ISI','IIIT-D'],'Rank':[0,1]},index=new_index)
      print(df1)
      print(df2)
```

```
      Name  Age
a  ankit   30
b   kiio   25
      College  Rank
a      ISI     0
b  IIIT-D     1
```

Index objects are immutable and thus can't be modified by the user

```
[46]: label=pd.Index(np.arange(2))
      df2=pd.DataFrame({'College':['ISI','IIIT-D'],'Rank':[0,1]},index=label)
      print(df2)
```

```
      College  Rank
0      ISI     0
1  IIIT-D     1
```

```
[47]: # reindexing

      df1=pd.DataFrame({'Name':['ankit','kiio'],'Age':[30,25]},index=['a','b'])
      df1=df1.reindex(['b','a','z'])
      print(df1)
```

```
      Name  Age
b   kiio  25.0
a  ankit  30.0
z    NaN   NaN
```

For ordered data like time series, it may be desirable to do some interpolation or filling of values when reindexing. The method option allows us to do this, using a method such as ffill, which forward-fills the values.

```
[48]: df1=pd.DataFrame({'Name':['ankit','kiio'],'Age':[30,25]},index=['a','b'])
      df1=df1.reindex(['b','a','z'],method='ffill')
```



```
print(df1)
```

```
      Name  Age
b    kiio   25
a   ankit   30
z    kiio   25
```

```
[49]: df=pd.DataFrame([[1.2,'ankit'],[9.7,'kiio'],[9.
↳5,'soumi']],index=['a','b','c'],columns=['Grade','Name'])
print(df)
df.drop('a',inplace=True)
print(df)
```

```
      Grade  Name
a      1.2  ankit
b      9.7  kiio
c      9.5  soumi
      Grade  Name
b      9.7  kiio
c      9.5  soumi
```

```
[50]: df.drop(['Grade'],axis=1)
```

```
[50]:      Name
b    kiio
c    soumi
```

```
[51]: df.drop(index=['b','c'])
```

```
[51]: Empty DataFrame
Columns: [Grade, Name]
Index: []
```

```
[52]: df=pd.Series(['ankit','kiio','summi','soumi'],index=['a','b','c','d'])
print(df[1:4])
print(df[['b','c']])
```

```
b    kiio
c    summi
d    soumi
dtype: object
b    kiio
c    summi
dtype: object
```

```
[53]: ##### NER (Named Entity Recognition) #####
```

```

import spacy

# Load the English model
nlp = spacy.load("en_core_web_sm")

text = "Barack Obama was born in Hawaii and was the president of the United_
↪States."

# Process the text
doc = nlp(text)

# Extract and print named entities
for ent in doc.ents:
    print(ent.text, ent.label_)

```

```

Barack Obama PERSON
Hawaii GPE
the United States GPE

```

```

[54]: df=pd.
      ↪DataFrame(['ankit','kiio','summi','soumi'],index=['a','b','c','d'],columns=['Name'])
print(df[1:4]) # selecting rows
print(df[['Name']]) # selecting columns
print(df[df['Name']=='kiio']) # filtering condition

```

```

      Name
b  kiio
c  summi
d  soumi
      Name
a  ankit
b  kiio
c  summi
d  soumi
      Name
b  kiio

```

```

[55]: df=pd.DataFrame([[1.2,'ankit'],[9.7,'kiio'],[9.
      ↪5,'soumi']],index=['a','b','c'],columns=['Grade','Name'])
print(df)
print(df.loc['a',['Grade','Name']])

```

```

      Grade  Name
a      1.2  ankit
b      9.7  kiio
c      9.5  soumi
Grade      1.2

```

```
Name      ankit
Name: a, dtype: object
```

```
[56]: print(df.iloc[1,[1,0]])
```

```
Name      kiio
Grade      9.7
Name: b, dtype: object
```

```
[57]: s1 = pd.Series([7.3, -2.5, 3.4, 1.5], index=['a', 'c', 'd', 'e'])
s1
```

```
[57]: a    7.3
      c   -2.5
      d    3.4
      e    1.5
      dtype: float64
```

```
[58]: s2 = pd.Series([-2.1, 3.6, -1.5, 4, 3.1], index=['a', 'c', 'e', 'f', 'g'])
s2
```

```
[58]: a   -2.1
      c    3.6
      e   -1.5
      f    4.0
      g    3.1
      dtype: float64
```

```
[59]: s1+s2
```

```
[59]: a    5.2
      c    1.1
      d   NaN
      e    0.0
      f   NaN
      g   NaN
      dtype: float64
```

```
[60]: df1 = pd.DataFrame(np.arange(12.).reshape((3, 4)), columns=list('abcd'))
df1
```

```
[60]:   a    b    c    d
0  0.0  1.0  2.0  3.0
1  4.0  5.0  6.0  7.0
2  8.0  9.0 10.0 11.0
```

```
[61]: df2 = pd.DataFrame(np.arange(20.).reshape((4, 5)), columns=list('abcde'))
df2
```

```
[61]:
```

	a	b	c	d	e
0	0.0	1.0	2.0	3.0	4.0
1	5.0	6.0	7.0	8.0	9.0
2	10.0	11.0	12.0	13.0	14.0
3	15.0	16.0	17.0	18.0	19.0

```
[62]: df2.loc[1, 'b'] = np.nan
```

```
[63]: df2
```

```
[63]:
```

	a	b	c	d	e
0	0.0	1.0	2.0	3.0	4.0
1	5.0	NaN	7.0	8.0	9.0
2	10.0	11.0	12.0	13.0	14.0
3	15.0	16.0	17.0	18.0	19.0

```
[64]: df1
```

```
[64]:
```

	a	b	c	d
0	0.0	1.0	2.0	3.0
1	4.0	5.0	6.0	7.0
2	8.0	9.0	10.0	11.0

```
[65]: df1 + df2
```

```
[65]:
```

	a	b	c	d	e
0	0.0	2.0	4.0	6.0	NaN
1	9.0	NaN	13.0	15.0	NaN
2	18.0	20.0	22.0	24.0	NaN
3	NaN	NaN	NaN	NaN	NaN

```
[66]: df1*df2
```

```
[66]:
```

	a	b	c	d	e
0	0.0	1.0	4.0	9.0	NaN
1	20.0	NaN	42.0	56.0	NaN
2	80.0	99.0	120.0	143.0	NaN
3	NaN	NaN	NaN	NaN	NaN

```
[67]: df1.add(df2, fill_value=0) # either in df1 or df2, wherever the value is nan,
    ↪ it will be replaced by zero and
    # then do df1+df2 as usual
```

```
[67]:
```

	a	b	c	d	e
0	0.0	2.0	4.0	6.0	4.0
1	9.0	5.0	13.0	15.0	9.0
2	18.0	20.0	22.0	24.0	14.0

```
3 15.0 16.0 17.0 18.0 19.0
```

```
[68]: arr = np.arange(12.).reshape((3, 4))
      arr
```

```
[68]: array([[ 0.,  1.,  2.,  3.],
            [ 4.,  5.,  6.,  7.],
            [ 8.,  9., 10., 11.]])
```

```
[69]: arr[0]
```

```
[69]: array([0., 1., 2., 3.]
```

```
[70]: arr-arr[0]
```

```
[70]: array([[0., 0., 0., 0.],
            [4., 4., 4., 4.],
            [8., 8., 8., 8.]])
```

When we subtract `arr[0]` from `arr`, the subtraction is performed once for each row. This is referred to as broadcasting.

If you want to instead broadcast over the columns, matching on the rows, you have to use one of the arithmetic methods.

```
[71]: frame=pd.DataFrame(np.arange(12).
      ↪reshape(4,3),columns=list('bde'),index=['Utah','Ohio','Texas','Oregon'])
      frame
```

```
[71]:
```

	b	d	e
Utah	0	1	2
Ohio	3	4	5
Texas	6	7	8
Oregon	9	10	11

```
[72]: series3 = frame['d']
      series3
```

```
[72]:
```

Utah	1
Ohio	4
Texas	7
Oregon	10

Name: d, dtype: int64

```
[73]: frame.sub(series3, axis='index')
```

```
[73]:
```

	b	d	e
Utah	-1	0	1

```
Ohio    -1  0  1
Texas   -1  0  1
Oregon  -1  0  1
```

The axis number that you pass is the axis to match on. In this case we mean to match on the DataFrame's row index (axis='index' or axis=0) and broadcast across.

```
[74]: # NumPy ufuncs (element-wise array methods) also work with pandas objects:

frame = pd.DataFrame(np.random.randn(4, 3), columns=list('bde'), index=['Utah', 'Ohio', 'Texas', 'Oregon'])
frame
```

```
[74]:
```

	b	d	e
Utah	0.829709	-0.077173	-0.072410
Ohio	-0.328474	-0.322371	-0.581475
Texas	0.619285	0.236105	-1.212218
Oregon	0.043921	1.539648	-0.300668

```
[75]: np.abs(frame)
```

```
[75]:
```

	b	d	e
Utah	0.829709	0.077173	0.072410
Ohio	0.328474	0.322371	0.581475
Texas	0.619285	0.236105	1.212218
Oregon	0.043921	1.539648	0.300668

Another frequent operation is applying a function on one-dimensional arrays to each column or row. DataFrame's apply method does exactly this.

```
[76]: frame
```

```
[76]:
```

	b	d	e
Utah	0.829709	-0.077173	-0.072410
Ohio	-0.328474	-0.322371	-0.581475
Texas	0.619285	0.236105	-1.212218
Oregon	0.043921	1.539648	-0.300668

```
[77]: f = lambda x: x.max() - x.min()
frame.apply(f)
```

```
[77]: b    1.158182
      d    1.862019
      e    1.139808
      dtype: float64
```

```
[78]: # If you pass axis='columns' to apply, the function will be invoked once per
      ↪row instead
      frame.apply(f, axis='columns')
```

```
[78]: Utah      0.906882
      Ohio      0.259104
      Texas     1.831503
      Oregon    1.840317
      dtype: float64
```

```
[79]: def f(x):
      return pd.Series([x.min(), x.max()], index=['min', 'max'])
      frame.apply(f)
```

```
[79]:          b          d          e
      min -0.328474 -0.322371 -1.212218
      max  0.829709  1.539648 -0.072410
```

```
[80]: def f(x):
      return pd.Series([x.min(), x.max()], index=['min', 'max'])
      frame.apply(f,axis='columns')
```

```
[80]:          min          max
      Utah  -0.077173  0.829709
      Ohio  -0.581475 -0.322371
      Texas -1.212218  0.619285
      Oregon -0.300668  1.539648
```

Element-wise Python functions can be used, too. Suppose you wanted to compute a formatted string from each floating-point value in frame. You can do this with apply map.

```
[81]: format = lambda x: '%.2f' % x
      frame.applymap(format)
```

/tmp/ipykernel_5585/1073433956.py:2: FutureWarning: DataFrame.applymap has been deprecated. Use DataFrame.map instead.

```
      frame.applymap(format)
```

```
[81]:          b          d          e
      Utah    0.83  -0.08  -0.07
      Ohio   -0.33  -0.32  -0.58
      Texas   0.62   0.24  -1.21
      Oregon  0.04   1.54  -0.30
```

```
[82]: frame['e'].map(format)
```

```
[82]: Utah      -0.07
      Ohio      -0.58
      Texas     -1.21
      Oregon    -0.30
      Name: e, dtype: object
```

```
[83]: pd.Series([1,3,2,7],index=['d','b','a','c']).sort_index()
```

```
[83]: a      2
      b      3
      c      7
      d      1
      dtype: int64
```

```
[84]: frame = pd.DataFrame(np.arange(8).reshape((2, 4)),index=['three',
↳ 'one'],columns=['d', 'a', 'b', 'c'])
      frame.sort_index()
```

```
[84]:      d  a  b  c
      one  4  5  6  7
      three 0  1  2  3
```

```
[85]: frame.sort_index(axis=1,ascending=True)
```

```
[85]:      a  b  c  d
      three 1  2  3  0
      one   5  6  7  4
```

```
[86]: obj = pd.Series([4, 7, -3, 2])
      obj.sort_values()
```

```
[86]: 2    -3
      3     2
      0     4
      1     7
      dtype: int64
```

```
[87]: frame = pd.DataFrame({'b': [4, 7, -3, 2], 'a': [0, 1, 0, 1]})
      frame
```

```
[87]:      b  a
      0  4  0
      1  7  1
      2 -3  0
      3  2  1
```

```
[88]: frame.sort_values(by='b')
```



```
[88]:    b  a
      2 -3  0
      3  2  1
      0  4  0
      1  7  1
```

```
[89]: frame.sort_values(by=['a', 'b']) # first sort by a and then b and when values_
      ↪ of a are same then for that
      # values of b will be sorted
```

```
[89]:    b  a
      2 -3  0
      0  4  0
      3  2  1
      1  7  1
```

```
[91]: df = pd.DataFrame([[1.4, np.nan], [7.1, -4.5], [np.nan, np.nan], [0.75, -1.
      ↪ 3]], index=['a', 'b', 'c', 'd'],\
      columns=['one', 'two'])
df
```

```
[91]:    one  two
a  1.40 NaN
b  7.10 -4.5
c   NaN NaN
d  0.75 -1.3
```

```
[92]: df.sum()
```

```
[92]: one    9.25
      two   -5.80
      dtype: float64
```

```
[93]: df.sum(axis=1)
```

```
[93]: a    1.40
      b    2.60
      c    0.00
      d   -0.55
      dtype: float64
```

Some methods, like `idxmin` and `idxmax`, return indirect statistics like the index value where the minimum or maximum values are attained.

```
[94]: df.describe()
```

```
[94]:
```

	one	two
count	3.000000	2.000000
mean	3.083333	-2.900000
std	3.493685	2.262742
min	0.750000	-4.500000
25%	1.075000	-3.700000
50%	1.400000	-2.900000
75%	4.250000	-2.100000
max	7.100000	-1.300000

```
[95]: df.idxmax()
```

```
[95]: one    b
      two    d
      dtype: object
```

```
[96]: df.idxmin()
```

```
[96]: one    d
      two    b
      dtype: object
```

Method → Description

count → Number of non-NA values
describe → Compute set of summary statistics for Series or each DataFrame column
min, max → Compute minimum and maximum values
argmin, argmax → Compute index locations (integers) at which minimum or maximum value obtained, respectively
idxmin, idxmax → Compute index labels at which minimum or maximum value obtained, respectively
quantile → Compute sample quantile ranging from 0 to 1
sum → Sum of values
mean → Mean of values
median → Arithmetic median (50% quantile) of values
mad → Mean absolute deviation from mean value
prod → Product of all values
var → Sample variance of values
std → Sample standard deviation of values
skew → Sample skewness (third moment) of values
kurt → Sample kurtosis (fourth moment) of values
cumsum → Cumulative sum of values
cummin, cummax → Cumulative minimum or maximum of values, respectively
cumprod → Cumulative product of values
diff → Compute first arithmetic difference (useful for time series)
pct_change → Compute percent changes

```
[97]: df.count()
```

```
[97]: one    3  
      two    2  
      dtype: int64
```

```
[98]: df.min()
```

```
[98]: one    0.75  
      two   -4.50  
      dtype: float64
```

```
[99]: df.max()
```

```
[99]: one    7.1  
      two   -1.3  
      dtype: float64
```

```
[100]: df.quantile()
```

```
[100]: one    1.4  
      two   -2.9  
      Name: 0.5, dtype: float64
```

```
[101]: df.sum()
```

```
[101]: one    9.25  
      two   -5.80  
      dtype: float64
```

```
[102]: df.mean()
```

```
[102]: one    3.083333  
      two   -2.900000  
      dtype: float64
```

```
[103]: df.median()
```

```
[103]: one    1.4  
      two   -2.9  
      dtype: float64
```

```
[104]: df.prod()
```

```
[104]: one    7.455  
      two    5.850  
      dtype: float64
```

```
[105]: df.var()
```

```
[105]: one    12.205833  
      two     5.120000  
      dtype: float64
```

```
[106]: df.std()
```

```
[106]: one     3.493685  
      two     2.262742  
      dtype: float64
```

```
[107]: df.skew()
```

```
[107]: one     1.664846  
      two      NaN  
      dtype: float64
```

```
[108]: df.kurt()
```

```
[108]: one      NaN  
      two      NaN  
      dtype: float64
```

```
[109]: df.cumsum()
```

```
[109]:      one  two  
a  1.40  NaN  
b  8.50 -4.5  
c   NaN  NaN  
d  9.25 -5.8
```

```
[110]: df.cumprod()
```

```
[110]:      one  two  
a  1.400  NaN  
b  9.940 -4.50  
c   NaN  NaN  
d  7.455  5.85
```

```
[111]: df.cummin()
```

```
[111]:      one  two  
a  1.40  NaN  
b  1.40 -4.5  
c   NaN  NaN  
d  0.75 -4.5
```

```
[112]: df.cummax()
```

```
[112]:      one  two
a   1.4  NaN
b   7.1 -4.5
c   NaN  NaN
d   7.1 -1.3
```

```
[113]: df.diff()
```

```
[113]:      one  two
a   NaN  NaN
b   5.7  NaN
c   NaN  NaN
d   NaN  NaN
```

```
[114]: df.pct_change()
```

```
/tmp/ipykernel_5585/890640361.py:1: FutureWarning: The default fill_method='pad'
in DataFrame.pct_change is deprecated and will be removed in a future version.
Either fill in any non-leading NA values prior to calling pct_change or specify
'fill_method=None' to not fill NA values.
df.pct_change()
```

```
[114]:      one      two
a      NaN      NaN
b  4.071429      NaN
c  0.000000  0.000000
d -0.894366 -0.711111
```

The corr method of Series computes the correlation of the overlapping, non-NA, aligned-by-index values in two Series. Relatedly, cov computes the covariance

```
[115]: df['one'].corr(df['two'])
```

```
[115]: np.float64(-1.0)
```

```
[116]: df['one'].cov(df['two'])
```

```
[116]: np.float64(-10.16)
```

value_counts -> Return a Series containing unique values as its index and frequencies as its values, ordered count in descending order

```
[117]: import pandas as pd
data = pd.DataFrame({'Qu1': [1, 3, 4, 3, 4], 'Qu2': [2, 3, 1, 2, 3], 'Qu3': [1, 4,
↪5, 2, 4, 4]})
data
```

```
[117]:
```

	Qu1	Qu2	Qu3
0	1	2	1
1	3	3	5
2	4	1	2
3	3	2	4
4	4	3	4

```
[118]: result = data.apply(pd.value_counts).fillna(0)
result
```

```
/tmp/ipykernel_5585/1382616601.py:1: FutureWarning: pandas.value_counts is
deprecated and will be removed in a future version. Use
pd.Series(obj).value_counts() instead.
  result = data.apply(pd.value_counts).fillna(0)
```

```
[118]:
```

	Qu1	Qu2	Qu3
1	1.0	1.0	1.0
2	0.0	2.0	1.0
3	2.0	2.0	0.0
4	2.0	0.0	2.0
5	0.0	0.0	1.0

```
[ ]:
```