

LLAMA INDEX

LlamaIndex

Source: Notes.md

LlamaIndex Udemy Course by Eden Marco

1. Large Language Models are powerful however they have some limitations.

For example, if I ask questions like **“How much did I pay in Serpapi in June ?”** to ChatGPT with GPT-3.5, it will not answer because it was not trained on this data but if I use GPT-4 and ask the question by attaching the invoice file, I will get the answer.

2. LlamaIndex

LlamaIndex is a powerful framework for developing LLM-based applications. It is open source and very intuitive to use. On GitHub, it has more popularity (based on the number of stars) than LangChain (a framework to develop LLM-based applications).

LlamaIndex is used to connect with our private data in:

- pdf
- ppt
- doc
- image
- audio
- video

as well as various data connectors API like:

- notion
- salesforce data
- discord

as well as various databases like:

- snowflake
- mongodb
- postgresql

LlamaIndex can also interact with vectorstores like **milvus** etc.

If we have a very large document then we have to do **chunking** and make various chunks because LLMs have token limits.

For a question, we might not need all the documents. We may get answer from:

- one specific chunk
- or 2-3 chunks

We only send these specific chunks as **context to LLM** then LLM will answer it.

We store these chunks into a **vector store/vector db** and retrieve relevant chunks based on the query/question.

3. Data Connectors (LlamaHub)

We have different data connectors (LlamaHub).

A **data connector (reader)** ingests data from different data sources and formats into a simple **Document** representation (text and simple data).

Once you have ingested your data, you can:

- build an **index**
 - ask questions using a **Query Engine**
 - have a conversation using a **Chat Engine**
-

4. LlamaHub

Our data connectors are offered through **LlamaHub**.

LlamaHub is an open source repository containing data loaders that you can easily plug and play into any LlamaIndex application.

5. Documents / Nodes

Document and Node objects are core abstractions within LlamaIndex.

Document

A **Document** is a generic container around any data source such as:

- PDF
- API output
- database data

They can be constructed manually or created automatically via data loaders.

A Document stores:

- text
 - metadata (dictionary of annotations)
 - relationships (dictionary linking other Documents/Nodes)
-

Node

A **Node** represents a **chunk of a source document** whether it is:

- text chunk
- image
- other data

Similar to Documents they contain metadata and relationship information to other nodes.

Node and chunk are same.

Nodes are a **first-class citizen in LlamaIndex**.

You can either:

- define Nodes directly
- parse documents into Nodes using **NodeParser**

By default every Node derived from a Document **inherits metadata** from that Document.

Example:

```
file_name
```

If it exists in Document metadata it will propagate to every Node.

6. Node Parser

Node parsers take a list of documents and **chunk them into Node objects**.

Each node has a **specific size**.

A node parser can configure:

- chunk size (tokens)
- overlap between chunked nodes

Chunking is done using:

TokenTextSplitter

Default values:

- chunk size = 1024
- chunk overlap = 20 tokens

7. Indexes

An **Index** is a data structure that allows us to quickly retrieve **relevant context** for a user query.

For LlamaIndex, it's the **core foundation for RAG use-cases**.

High-level structure:

Documents → **Index** → Query Engine / Chat Engine

Indices are built from **Documents**.

They are used to build:

- Query Engines
- Chat Engines

which enables question answering and chat over your data.

Under the hood:

Indices store data in **Node** objects and expose a **Retrieval interface** supporting automation and configuration.

8. Query Engine

Query engine is a generic interface that allows you to **ask questions over your data**.

A query engine:

- takes a natural language query
- returns a rich response

It is often built on one or many indices via Retrievals.

Multiple query engines can be composed for advanced capabilities.

9. LlamaIndex Overview

LlamaIndex is the leading framework for building LLM-powered agents over your data with LLMs and workflows.

Sections include:

- Introduction
 - Use cases
 - Getting started
 - LlamaCloud
 - Community
 - Related projects
-

Introduction

What are agents?

Agents are LLM-powered knowledge assistants that use tools to perform tasks like:

- research
- data extraction
- task execution

Agents range from simple question answering to systems that **sense, decide and take actions**.

LlamaIndex provides a framework for building agents including the ability to use **RAG pipelines as tools**.

What are workflows?

Workflows are multi-step processes that combine:

- agents
- data connectors
- tools

They are **event-driven systems** allowing complex applications with:

- reflection
- error correction
- reasoning

These workflows can be deployed as **production microservices**.

What is context augmentation?

LLMs are trained on large public datasets but not on **your private data**.

Your data may exist in:

- APIs
- SQL databases
- PDFs
- slide decks

Context augmentation makes your data available to the LLM.

The most popular example is:

Retrieval Augmented Generation (RAG)

which combines **context** + **LLM at inference time**.

LlamaIndex provides tools such as

Data connectors

Data indexes

Query engines

Chat engines

Agents

Observability tools

Workflows

Use Cases

Popular use cases include:

- Question answering (RAG)
 - Chatbots
 - Document understanding
 - Data extraction
 - Autonomous agents
 - Multimodal applications
 - Model fine tuning
-

Who is LlamaIndex for?

LlamaIndex provides tools for:

Beginners

Use high level APIs to ingest and query data in **5 lines of code**

Advanced users

Customize modules like:

- data connectors
- indices
- retrievers
- query engines
- reranking modules

Getting Started

LlamaIndex is available in:

- Python
- TypeScript

30 second quickstart

Set environment variable:

```
OPENAI_API_KEY
```

Install library:

```
pip install llama-index
```

BASH

Example code:

```
from llama_index.core import VectorStoreIndex, SimpleDirectoryReader

documents = SimpleDirectoryReader("data").load_data()
index = VectorStoreIndex.from_documents(documents)
query_engine = index.as_query_engine()

response = query_engine.query("Some question about the data should go here")

print(response)
```

PYTHON

LlamaCloud

LlamaCloud is a managed service for:

- document parsing
- data extraction
- indexing
- retrieval

Website:

<https://llamaindex.ai/enterprise>

You can get **10,000 free credits per month**.

Services include:

Document Parsing (LlamaParse)

Document Extraction (LlamaExtract)

Indexing / Retrieval pipelines

Community

Twitter

Discord

LinkedIn

Library Links

Python:

- https://github.com/run-llama/llama_index
- <https://docs.llamaindex.ai/>
- <https://pypi.org/project/llama-index/>

TypeScript:

- <https://github.com/run-llama/LlamaIndexTS>
- <https://ts.llamaindex.ai/>
- <https://www.npmjs.com/package/llamaindex>

10. Documentation

https://github.com/run-llama/llama_index

python-dotenv:

<https://github.com/theskumar/python-dotenv>

https://www.youtube.com/watch?v=8dlQ_nDE7dQ&t=210s

11. Project Setup

Use **Python 3.11** since it's the most stable for LlamaIndex.

Create project directory:

```
mkdir llamaindex-helloworld  
cd llamaindex-helloworld
```

BASH

Install pipenv:

```
pip3 install pipenv  
pipenv --version
```

BASH

Create virtual environment:

```
pipenv shell
```

BASH

Exit environment:

```
exit
```

If creating a fresh environment:

```
rm -f Pipfile Pipfile.lock  
  
export PIPENV_NO_INHERIT=true  
export PIPENV_IGNORE_REQUIREMENTS=true  
  
pipenv --python 3.11  
  
pipenv install llama-index python-dotenv  
  
pipenv shell
```

BASH

Install packages:

BASH

```
pipenv install llama-index python-dotenv  
pipenv install llama-index-readers-web html2text
```

Create files:

BASH

```
touch main.py .env
```

Run program:

BASH

```
pipenv run python main.py
```

12. Retrieval Augmented Generation (RAG)

Prompt = **Query** + **Context**

Large text (like books) must be **chunked** to avoid token limits.

Only relevant chunks are sent to the LLM.

This reduces unnecessary API costs.

Embeddings

Embeddings convert objects into **vectors**.

Objects may include:

- text
- audio
- image
- video

Similar semantic meanings produce **vectors close together in embedding space**.

Images:

EMBEDDINGS.

Query

"How tall is the Burj Khalifa?"

Context

The Burj Khalifa, known as the Burj Dubai prior to its inauguration in 2010, is a skyscraper in Dubai, United Arab Emirates. It is known for being the world's tallest building. With a total height of 829.8 m (2,722 ft, or just over half a mile) and a roof height (excluding antenna, but including a 242.6 m spire[2]) of 828 m (2,717 ft)

Encoder/
embedding model

vector space

Query

Context

PROMPT AUGMENTATION

Prompt

Query = 'What is the burj khalifa ?'

+Context =

13. Language Modelling

A language model is a probability distribution over sequences of words.

Formal definition:

Given sequence

$x_1, x_2, x_3 \dots x_t$

Predict:

$P(x_{t+1} \mid x_1, x_2, \dots, x_t)$

Examples:

- search engines
- text suggestions

14. Large Language Model (LLM)

An LLM is simply a **language model trained on huge amounts of data**.

These models are very good at calculating probabilities.


15. Composition of a Prompt

A prompt is the **input given to a language model**.

Prompt components:

- A. Instructions
- B. Context
- C. Input Data
- D. Output Indicator

Example diagram:



Prompt

16. Prompting Techniques

Zero Shot Prompting

Model generates output without examples.

Example:

Write an image description with adjectives and nouns of a Yorkshire dog running in a winter landscape in Brazil.

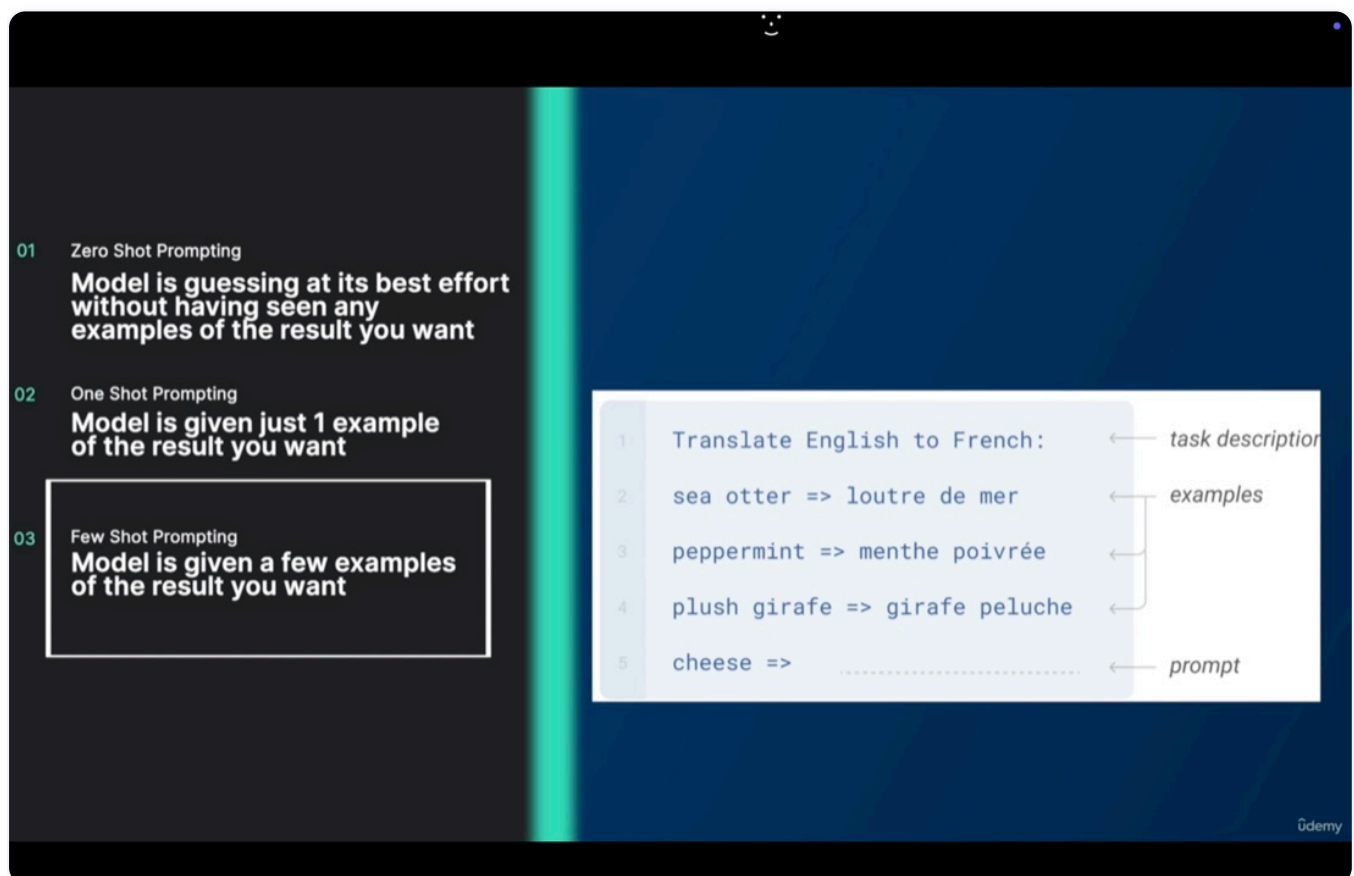
One Shot Prompting

Model receives **one example**.

Few Shot Prompting

Model receives **multiple examples**, improving accuracy.

Image:



Few Shot Prompt

One shot prompting is a subset of few shot prompting.

17. Chain of Thought Prompting

Large LLMs still struggle with **multi-step reasoning problems**.

In 2022 Google researchers introduced **Chain of Thought (CoT)**.

It improves reasoning by breaking problems into **intermediate steps**.

Zero Shot CoT:

Prefix prompt with:

```
Let's think step by step.
```

Example:

```
I bought 10 apples.  
Gave 2 to neighbor.  
Gave 2 to repairman.  
Bought 5 more.  
Ate 1.
```

```
How many apples remain?
```

```
Let's think step by step.
```

Few Shot CoT

Provide examples of **step-by-step reasoning**.

Images:

Standard prompting

Input: Q: Shawn has five toys. For Christmas, he got two toys each from his mom and dad. How many toys does he have now?
A: The answer is 9.

...

Q: John takes care of 10 dogs. Each dog takes .5 hours a day to walk and take care of their business. How many hours a week does he spend taking care of dogs?

A:

Model output: The answer is 50. ❌

Chain of thought prompting

Gemini

Chain of thought prompting

Input: Q: Shawn has five toys. For Christmas, he got two toys each from his mom and dad. How many toys does he have now?
A: Shawn started with 5 toys. If he got 2 toys each from his mom and dad, then that is 4 more toys. $5 + 4 = 9$. The answer is 9.

...

Q: John takes care of 10 dogs. Each dog takes .5 hours a day to walk and take care of their business. How many hours a week does he spend taking care of dogs?

A:

Model output: John takes care of 10 dogs. Each dog takes .5 hours a day to walk and take care of their business. So that is $10 \times .5 = 5$ hours a day. $5 \text{ hours a day} \times 7 \text{ days a week} = 35 \text{ hours a week}$. The answer is 35 hours a week. ✅

Gemini

18. ReAct Prompt Engineering

ReAct stands for:

Reason + Act

ReAct combines:

CoT + Actions

It integrates reasoning with actions and enables interaction with external environments.

It comes from a **2023 research paper** and forms the basis of **LangChain agents**.

19.